

The Tattoo as Narrative Archive: a situated reading of digital sources with a local system

Abstract

The tattoo can be read as a bodily archive, an aesthetic mark, an affective memory, a visual craft, and an object of social dispute. Digital sources, however, create a methodological difficulty: collecting news, forums, reports, and academic articles may produce a large corpus while turning heterogeneous voices into a falsely unified discourse. This paper presents a locally grounded procedure for building a traceable narrative corpus without sending texts to external services. The guiding case is *tattoo*, treated not as a single keyword but as a field of meanings: body, identity, craft, aesthetics, health, regulation, memory, and stigma. The system separates search rubrics, years, and source layers; cleans texts; records gaps and failures; helps organize narrative signals, build maps of relations, and propose revisable syntheses for close reading. The aim is not to automate interpretation, but to provide an auditable methodological mediation for historians, artists, designers, and cultural studies scholars.

1 Cultural and methodological problem

A social narrative does not appear in only one kind of document. It may emerge in everyday conversation, circulate through journalism, be translated into institutional language, and later become stabilized in academic research. In the case of tattooing, the same term may refer to artistic practices, commercial brands, medical procedures, youth identities, workplace stigma, or sanitary regulation. If everything is collected together, the corpus may grow while its interpretive value decreases.

The purpose of the procedure is to build a situated corpus that preserves three things: provenance, difference among voices, and cleaning decisions. The question is not “what does the Internet say about tattoos?” but how different discursive layers produce, move, or dispute meanings around tattooing in a defined period and region.

2 Interpretive phases

Phase	What is done	Why it matters
Situating the field	Tattooing is recognized as bodily inscription, visual practice, and social narrative.	The cultural problem is framed before it becomes a digital search.
Delimiting	Region, period, and meaning rubrics are defined: identity, craft, health, memory, design.	Different meanings are not collapsed under one term.
Gathering voices	Sources are collected by layer: press, forums, articles, and reports.	Public conversation, journalism, research, and institutional documents remain distinguishable.
Caring for the corpus	Duplicates, menus, advertising, noise, and homonyms are removed.	Web noise is prevented from becoming interpretive evidence.
Reading relations	Expressions, actors, narrative moments, sources, and lexical tonality are observed.	The system produces clues and maps of meaning for critical reading.
Critical synthesis	Relevant nodes are proposed and checked against textual examples.	The algorithm is not treated as the final interpretation.

Table 1: Methodological phases for a situated reading of tattooing in digital sources.




Figure 1: Interpretive itinerary: the system organizes reading; it does not replace the researcher.

3 Guiding case: tattooing

For historians, artists, designers, and cultural studies scholars, tattooing should not be treated as an isolated keyword. It should be divided into interpretive rubrics:

Rubric	Search variants
Core terms	tattoo, tattoos, tatuaje, tatuajes, body art
Craft and industry	tattoo artist, tattoo studio, tatuador, tatuadora, artist of tattooing
Identity and memory	tattoo meaning, tattoo and identity, tattoo and memory, religious tattoo
Society and work	tattoo and youth, tattoo and employment, tattoo discrimination, tattoo and gender
Health and regulation	tattoo health, tattoo inks, infection, sanitary regulation
Aesthetics and design	tattoo design, traditional tattoo, Mexican tattoo, body ink

Table 2: Initial rubrics for the tattoo case. Rubrics are editable and must be validated through reading.

Exclusions are also declared. They are necessary because *Tatuaje* may appear as a cigar brand or tattooing may refer to specialized medical procedures. Excluding terms such as *cigar*, *tobacco*, *robusto*, or *colonoscopic tattooing* does not censor the corpus; it protects the research question.




Figure 2: Corpus design for the tattoo case: notice how the cultural question, meaning rubrics, and voice layers are separated before documents are collected.

4 From documents to narrative traces

The system saves a separate base for each combination of rubric, year, and source layer. If a search produces no results, the absence is registered. This matters: an absence also speaks about the scope of the corpus and about the limits of collection.

The current collection strategy is sequential and adaptive: it queries the base term of a rubric first and expands to synonyms only when more evidence is needed. This reduces pressure on public indexes and preserves the term that activated each finding. In the tattoo case, collecting through *tatuaje*, *tatuajes*, *tattoo*, or *body art* does not produce the same corpus; each term carries different communities, languages, and biases.

Step	Example	Methodological reading
Input	“tattoo and health”, Mexico, 2021, news	One sense of the topic is located in a public layer.
Cleaning	Menus, ads, duplicates, and homonyms are removed	The corpus does not confuse web noise with social narrative.
Description	Terms such as ink, skin, regulation, risk appear	These are clues, not conclusions.
Relation	Text–actor–idea–source– year	The system shows how voices and concepts are connected.
Synthesis	Relevant node: sanitary regulation / tattoo studio	The node becomes an entry point for close reading.

Table 3: Mini-demonstration of the procedure using the health and regulation rubric.

5 Source rights and degrees of evidence

Extraction must respect the limits of each source. The system does not enter private spaces, does not automate logged-in sessions, and does not bypass access restrictions. If a public page allows extraction according to `robots.txt`, the system attempts to obtain full text. If extraction is not allowed, if the site only exposes partial material, or if the content is behind a paywall, the system keeps only a traceable signal: title, snippet, outlet, date, URL, and source type. That signal may orient the corpus map, but it must not be counted as equivalent to full textual evidence.




Figure 3: Degrees of evidence: full text and partial signals are stored separately so the corpus does not exaggerate its evidentiary reach.

This distinction is central when studying what people say online. Forums, public Reddit pages, blogs, and open communities provide traces of public conversation, but they do not represent the whole social conversation. They must be read as indexable traces shaped by platform bias, language, moderation, visibility, and access.

Micro-case for reading

A publishable run should not stop at document counts. It should show at least one minimal chain: source, cleaned fragment, assigned label, constructed relation, and human reading. In the tattoo case, a news item about inks or sanitary regulation may connect health authorities, tattoo studios, and risk vocabulary; a forum conversation may connect clients, pain, identity, price, or trust; an academic article may stabilize medical vocabulary around skin, infection, or healing.

Layer	Possible signal	Node/relation	Human reading
Press	“ink regulation”	health authority–ink–risk	Tattooing appears as a public and sanitary issue.
Forum	“studio recommendation”	client–tattoo artist–trust	Craft is evaluated through shared experience.
Article	“skin infection”	skin–procedure–care	The body becomes an object of medical knowledge.

Table 4: Example of how a document becomes an entry for reading, not an automatic conclusion.

6 Maps of relations

The narrative network is a map. Each text connects actors, ideas, source, year, place, and moments of the story. The map does not prove a hypothesis by itself; it helps decide where close reading should begin.

Actors, ideas, and narrative moments should not be treated as transparent entities. They may come from local rules, editable dictionaries, n-gram frequencies, linguistic patterns, and human validation. In a topic such as tattooing, expected actors may include tattoo artists, clients, health authorities, physicians, employers, media outlets, and youth collectives; ideas may group identity, memory, craft, risk, aesthetics, discrimination, or care. Each assignment must remain correctable through reading.




Figure 4: Minimal narrative network: the reader can see what kind of relation is being recorded.

7 Exploratory lexical tonality

Lexical tonality is not interpreted as collective emotion or as the psychological state of authors. It is used as an exploratory indicator of evaluative vocabulary. It may show, for example, whether news sources describe tattooing through risk, whether forums use language of identity, or whether

academic articles concentrate on health and regulation.

The basic indicator compares positive and negative words, with simple rules for negation and intensification. Results are summarized by year and source layer. In a radial chart, each axis represents one layer: press, forums, articles, or reports. Interpretation must always return to textual examples.

8 Revisable synthesis

When the map contains too many elements, the system proposes relevant nodes: actors, ideas, sources, or narrative moments that concentrate relations. This selection does not replace interpretation. It helps ask which entries deserve close reading and what relations would be lost if certain nodes were removed.




Figure 5: Synthesis is revisable: the algorithm suggests entry points; interpretation returns to the corpus.

9 Reproducible output

The local application exports a structured evidence file that preserves filtered records, provenance, cleaning decisions, source classification, lexical tonality, narrative events, actors, idea groups, and maps of relations. For a stronger publication workflow, a run manifest should also be consolidated with date, parameters, code version, quotas, seeds, and a fingerprint of the final corpus.

10 Limits and biases

The corpus does not represent a society by itself. It represents a situated collection of available traces: what indexes find, what platforms allow one to consult, what media outlets publish, and what certain groups write online. It may contain biases of language, social class, digital literacy, urban centrality, source availability, platform moderation, journalistic agenda, and sanitary or institutional overrepresentation. Exclusions are also interpretive decisions: removing *colonoscopic tattooing* is reasonable if tattooing is studied as cultural practice, but it may be wrong if the rubric is medicine or health. Every exclusion should therefore be registered and justified.

A Minimal technical note

The selection of relevant nodes can be formalized as a coverage problem over the map of relations. The system searches for small sets of nodes that maintain high interpretive relevance and low structural damage. Operationally, three criteria are compared: set size, narrative weight of the nodes, and relations that would be lost if those nodes were removed. The algorithmic details are left in the appendix because the main question is interpretive: which set of actors, ideas, or sources offers a defensible entry point for reading the corpus?

B Update: available evidence and analysis requirements (October 6, 2026)

This update supersedes earlier assumptions of complete metadata availability. Authorship, publication dates and outlets are sought but cannot be guaranteed. Each field preserves its value, evidence, method and state: explicit, inferred, searched but not found, not evaluable, not applicable or conflicting. Conflicting candidates are retained rather than arbitrarily resolved.

Publication, update and consultation dates are distinct. Precision is year, month or day; missing components are never filled. The legacy search year is preserved as `search_year`; `year` only represents usable explicit publication evidence. Index observation dates are not publication dates. A PDF link alone is not evidence of full-text retrieval.

Documents without authors or dates remain available for thematic exploration. Each analysis reports eligible documents, corpus total, exclusion reasons and coverage by source type. Coverage does not establish representativeness or statistical sufficiency. Legacy ambiguous dates require human review for temporal comparison. A URL domain is not an author, identified publisher or event location.

B.1 Descriptive models and review

The document–concept matrix uses 1 for reviewed presence, 0 only for absence after explicit concept review, and missing values for unreviewed cases. Source–concept proportions use concept-specific reviewed denominators. Actor–concept–stance counts keep support, rejection and ambivalence separate, requiring a verbatim quotation and reviewed speaker attribution. Directed relationships retain their type, document and supporting quotation. Causation asserted in a text is not an estimated causal effect. Temporal counts cover the intersection of reviewed and explicitly dated documents, with coverage reported. They do not estimate population attitudes.

Existing heuristic modules produce exploratory candidates. `evidence_model.py` checks metadata states and quotations, not disciplinary truth. JSON reviews are preserved through corpus export. Reception must be supported by response evidence, not inferred from a message alone. Documents can depend on the same institution or event; inferential testing and diffusion estimation are not implemented by these descriptive counts.

B.2 Pseudocode

1. Retrieve content and provenance; search labelled metadata when possible.
2. Preserve field values, evidence, states, conflicts and date precision.
3. Separate publication, update, consultation and search year.
4. Check requirements per analysis; record inclusion or exclusion reason.
5. Check verbatim quotations and attribution in submitted reviews.
6. Build matrices with missing values distinct from zero; preserve stance categories.
7. Export coverage, matrices, relations and limitations of the used subset.

The complete policy and reproducible commands are in `EVIDENCIA_Y_MODELOS.md`. Updating the software does not turn historical reports into new measurements.

B.3 Shared corpus contract

Identity and evidence-preserving merging are centralized in `corpus_contract.py`, shared by the interface and command line tools. Identifier URL parameters, retrieval provenance, text variants and conflicts are retained. Similar titles do not establish document identity. Lexical non-detection in recovered fragments does not prove intentional silence. Annual dates do not establish within-year order. Relational exports retain empty reviewed-model tables when evidence is insufficient. The integral audit documents pilot limits.

B.4 Versioned records and stages

Document, version and retrieval identities are separate. Reviewed selection requires rationale, evidence and reviewer; changes invalidate prior reviews. Titles, abstracts and fragments remain distinct. Six stages retain integrity checked checkpoints for resumption. Saved-corpus processing organizes existing evidence rather than simulating new retrievals. Query caching, bounded retries, transactional storage and atomic writes support the versioned corpus contract.